

# Stat 243

## Problem set 6

A. Strand, ID: 540441, GitHub: AndreasStrand

November 2, 2015

### Problem 1

The code for problem 1 are included in code 1. The size of the original CSV is 12 Gb, the bzipped file is 1.7 Gb while the SQLite database is of size 9.4 Gb. The SQLite storage is in other words more efficient than the .csv format, but less efficient than the bz2-zipped csv.

Listing 1: R code for problem 1 with comments.

```
setwd("~/Documents/stat243")

# Running this code bunch in the EC2
library(RSQLite)
year = seq(1987,2008, by =1)
airNames = paste("http://www.stat.berkeley.edu/share/paciorek/",
year, ".csv.bz2", sep = "")
filename = "airline.db"
db <- dbConnect(dbDriver("SQLite"), dbname = filename)
for (i in 1:length(year)){
  dbWriteTable(conn = db, name = "airlines",
    value = read.csv(bzfile(airNames[i]), header=T, sep = ","),
    row.names = FALSE, header = TRUE, append = T)
}
# Disconnecting using
dbDisconnect(db)
```

### Problem 2

The answer for problem 2 is written as bash code with comments and pseudo code in code 2. Besides experimentation with the cluster, there is for some parts pseudo code to express the general idea. There was some error in my SQLite calculations in R, but for the speed comparison I would use system.time() in R and compared it to the equivalent for the pyspark RDD computations.

Listing 2: Bash code, pyspark code and pseudo code for problem 2 with comments.

```
setwd("~/Documents/stat243")
#### Logon to spark ####
export NUMBER_OF_WORKERS=12
cd ~/Documents/stat243/spark-1.5.1/ec2
export AWS_ACCESS_KEY_ID='grep aws_access_key_id ~/Documents/stat243/
stat243-fall-2015-credentials.boto | cut -d' ' -f3'
export AWS_SECRET_ACCESS_KEY='grep aws_secret_access_key ~/Documents/stat243/
stat243-fall-2015-credentials.boto | cut -d' ' -f3'
chmod 400 ~/Documents/stat243/.ssh/stat243-fall-2015-ssh_key.pem

./spark-ec2 -k andrstra@stud.ntnu.no:stat243-fall-2015 -i ~/Documents/
```

```

stat243 /.ssh/stat243-fall-2015-ssh_key.pem --region=us-west-2 -s
${NUMBER_OF_WORKERS} -v 1.5.1 launch sparkvm-andrstra@stud.ntnu.no
./spark-ec2 -k andrstra@stud.ntnu.no:stat243-fall-2015 -i ~/Documents
/stat243 /.ssh/stat243-fall-2015-ssh_key.pem --region=us-west-2 login sparkvm-
andrstra@stud.ntnu.no
# terminating later using
./spark-ec2 --region=us-west-2 --delete-groups
destroy sparkvm-andrstra@stud.ntnu.no

```

```
##### In spark #####
```

```
##### This section is containing pseudo code for problem 2
```

```
# Columns of airline data
```

[1]	"Year"	"Month"	"DayofMonth"
[4]	"DayOfWeek"	"DepTime"	"CRSDepTime"
[7]	"ArrTime"	"CRSArrTime"	"UniqueCarrier"
[10]	"FlightNum"	"TailNum"	"ActualElapsedTime"
[13]	"CRSElapsedTime"	"AirTime"	"ArrDelay"
[16]	"DepDelay"	"Origin"	"Dest"
[19]	"Distance"	"TaxiIn"	"TaxiOut"
[22]	"Cancelled"	"CancellationCode"	"Diverted"
[25]	"CarrierDelay"	"WeatherDelay"	"NASDelay"
[28]	"SecurityDelay"	"LateAircraftDelay"	

```
## Creating directories
```

```

export PATH=$PATH:/root/ephemeral-hdfs/bin/
hadoop fs -mkdir /data
hadoop fs -mkdir /data/airline
df -h
mkdir /mnt/airline

```

```
## Loading the airline data onto an HDFS
```

```

for i in `seq 1987 2008`;
do
    wget "http://www.stat.berkeley.edu/share/paciorek/${i}.csv.bz2"
done
hadoop fs -copyFromLocal /mnt/airline/*bz2 /data/airline
# Check files on the HDFS
hadoop fs -ls /data/airline

```

```
## Installing numpy and running pyspark
```

```

yum install -y python27-pip python27-devel
pip -2.7 install 'numpy==1.9.2'
export PATH=$PATH:/root/spark/bin
pyspark
from operator import add
import numpy as np

```

```
## Creating an RDD "lines"
```

```

lines = sc.textFile('/data/airline')
numLines = lines.count()
def stratify(line)
    vals = line.split(',')
    return (vals[16], 1)
result = lines.map(stratify).reduceByKey(add).collect()

```

```

## Filter subset so it does not contain missing or unreasonable values for
# departure delay. Also repartitioning data such that it is equally spread across
# the twelve worker nodes
lines.filter(lambda line: (~!NA >= 0 and < 300) in line.split(' ')[16]).repartition(12)

## Aggregating information to categories
# airline, departure airport, arrival airport, calendar month, day of week,
# hour of day of scheduled dep. time
# column numbers: 11, 17, 18, 2, 4, 5
# (for the last element take floor(<val 5>/100) to get the hour of the sch. dep. time)
def computeKeyValue(line):
    vals = line.split(' ')
    keyArr = [vals[x] for x in [11, 17, 18, 2, 4, 5]]
    keyArr[5] = floor(keyArr[5]/100)
    ## Creating key by joining column values by hyphen
    keyVals = '-'.join(keyArr)
    delay30 = (vals[16] >= 30)
    delay60 = (vals[16] >= 60)
    delay180 = (vals[16] >= 180)
    return (keyVals, delay30, delay60, delay180)

# I am not sure about details in reduce by key, but the plan is that since
# delay30, delay60 and delay180 is boolean the average is the same as the share.
# Furthermore if the reduction for each key is recursive by the structure
# mean(mean(mean(a), b), c) I think it could work with the lambda suggested
avlambdas = lambda v1, v2: numpy.average(v1, v2)
newRDD = lines.map(computeKeyValue).reduceByKey(avlambdas)

## Saving the aggregated data set onto the master node
def strat(line):
    changedLine = *change to comma-delimited char. string*
    return (changedLine)
newRDD.lines.map(strat).groupByKey().saveAsTextFile('/data/delayOverview')

```

## Problem 3

The `apply()` family of functions in R are really useful, and I am looking forward to checking out the parallel apply calculations. I will look into it as we use in real work flows in class. I have to admit that I did not have the time now. By the way, what is the index from *Pr.2c*?

## Problem 4

To extract certain columns from bash you may use code like the one below.

```
awk -F "\"*,\"" '{print <column numbers>}' textfile.csv
```

This operation is fast, but even if it is done using piping (`unzip|*awk*|zip`) the zipping and unzipping will take some time. I do not think this is a worthwhile step.